

Chaouki_Ayman_TP1

November 12, 2018

1 Dynamic Programming


1.0.1 Q1: Implementing the discrete MDP model

There are no rewards in state s_1 , therefore we want to avoid being stuck in it. We have $r(s_2, a_2) = \frac{9}{10}$, which is a good reward, thus we want to get stuck in state s_2 , thus we take:


$$\begin{cases} \pi^*(s_0) = a_1 \\ \pi^*(s_1) = a_1 \\ \pi^*(s_2) = a_2 \end{cases}$$

1.0.2 Q2 : Value iteration

We want a 0.01-optimal policy, therefore we take $\frac{2\epsilon\gamma}{1-\gamma} = 0.01$ where ϵ is the stopping criterion of the algorithm.



1.0.3 Q3 : Policy iteration



Timing value iteration and computing the number of iterations it takes until convergence.

6.78 ms \pm 1.18 ms per loop (mean \pm std. dev. of 7 runs, 100 loops each)

number of iterations taken by value iteration : 160

Timing policy iteration and computing the number of iterations it takes until convergence.

370 μ s \pm 49 μ s per loop (mean \pm std. dev. of 7 runs, 1000 loops each)

number of iterations taken by policy iteration : 4

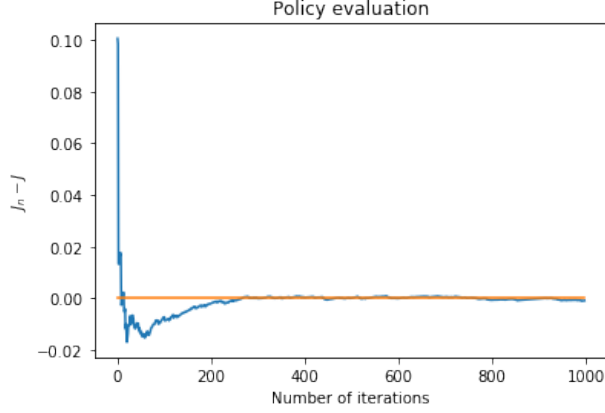
In this example, policy iteration is much faster to converge than value iteration and it takes much less iterations. The problem with policy iteration is that it requires a full policy evaluation in each step, thus if the state space is huge this step becomes expensive and value iteration may become the preferable option.

2 Reinforcement Learning

2.0.1 Q4 : Policy evaluation

We take $T_{max} = 140$ to get an approximation of around 10^{-3} and for each state we simulate 1000 trajectories to estimate the optimal value function.

```
Out[13]: array([ 0.87511567,  0.92961951,  0.98980574,  0.          ,  0.66599418,
                -0.99451987,  0.          , -0.82820401, -0.87698198, -0.93385322,
                -0.994725   ])
```



2.0.2 Q5 : Policy optimization

Assume we have some state-action function $Q_t(x, a)$, and we are in the current state $x_t \in \mathcal{X}$. We choose the next state x_{t+1} according to an ϵ -greedy policy, that is we choose $a_t = \operatorname{argmax}_{a \in \mathcal{A}} Q_t(x_t, a)$ with probability $1 - \epsilon$ otherwise we choose an action from the rest with a uniform distribution. Now that we have simulated the transition (x_t, a_t, r_t, x_{t+1}) , we can compute the temporal difference:

$$\delta_t = r_t + \gamma \cdot \max_{a \in \mathcal{A}} (Q_t(x_{t+1}, a)) - Q_t(x_t, a_t) \quad (1)$$

Then we can update the estimate:

$$Q_{t+1}(x_t, a_t) = Q_t(x_t, a_t) + \alpha_{N(x_t, a_t)} \delta_t \quad (2)$$


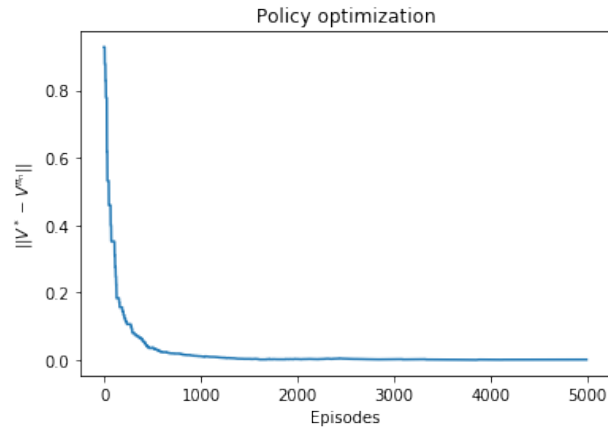
Where $N(x_t, a_t)$ is the number of times we visited state-action (x_t, a_t) , by taking $\alpha_n = \frac{1}{n^\beta}$ where β is a decay coefficient in $]0.5, 1]$, the learning rate satisfies Robbins-Monro conditions:

$$\begin{cases} \sum_n \alpha_n = \infty \\ \sum_n \alpha_n^2 < \infty \end{cases}$$

but in order for the Q-learning algorithm to converge, all states-actions must be visited infinitely often. This is relying on the quality of the exploratory policy. In this case, a good choice of ϵ .

Note: We may need to decrease ϵ to 0 over time with a cooling schedule like SARSA algorithm does with its exploratory policy.

The algorithm we will run uses $\epsilon = 0.3$ with a cooling schedule $\epsilon(x, a) = \frac{1}{\epsilon + N(x, a)}$ where $N(x, a)$ is the number of visits to the state-action (x, a) and in the same fashion $\alpha(x, a) = \frac{1}{N(x, a)^\beta}$ with $\beta = 0.8$



Q-learning algorithm does converge well to the optimal value function and optimal policy.

2.0.3 Q6:

The initial distribution μ_0 does influence the estimation of the optimal policy, **But not the true optimal policy since it is intrinsic to the MDP.**

- If μ_0 always gives an absorbant state, no trajectories will be produced and thus Robbins-Monro condition of visiting each state-action infinitely often does not hold.
- If the MDP simulator is deterministic with respect to some state-action, if μ_0 is poorly chosen, there is a risk of not exploring all the states. There will be a need for a very good exploratory policy, therefore μ_0 does influence the estimation of the optimal policy.